

# DL LAB1 Report - Back propagation

匯出pdf之後排版變得有點醜 QQ

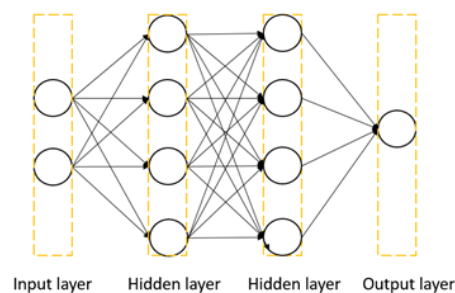
助教不介意的話可以到HackMD網址閱讀 ~ ~ ~ 謝謝

<https://hackmd.io/UsGKuF5yTVmfr3PGU0BKsw>

(<https://hackmd.io/UsGKuF5yTVmfr3PGU0BKsw>)

## 1. Introduction

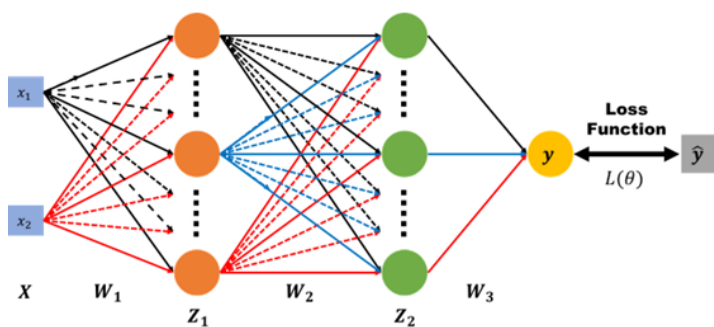
In this lab, we need to understand and implement simple neural networks with forwarding pass and backpropagation.



## Request

- Implement simple neural networks with two hidden layers.
- You must use backpropagation in this neural network and can only use Numpy and other python standard libraries to implement.
- Plot your comparison figure that show the predicted results and the ground-truth.

## Details



1.  $x_1, x_2$  : neural network inputs<sup>↵</sup>
2.  $X : [x_1, x_2]^{\epsilon \times 1}$
3.  $y$  : neural network outputs<sup>↵</sup>
4.  $\hat{y}$  : ground truth<sup>↵</sup>
5.  $L(\theta)$  : loss function<sup>↵</sup>
6.  $W_1, W_2, W_3$  : weight matrix of network layers

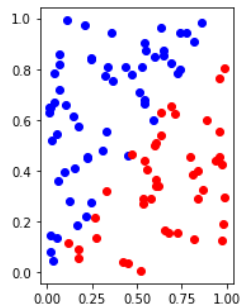
$$Z_1 = \sigma(XW_1) \quad Z_2 = \sigma(Z_1W_2) \quad y = \sigma(Z_2W_3)$$

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

# Dataset

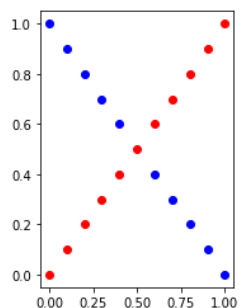
- Linear

```
1 def generate_linear(n=100):
2     pts = np.random.uniform(0, 1, (n, 2))
3     inputs = []
4     labels = []
5     for pt in pts:
6         inputs.append([pt[0], pt[1]])
7         distance = (pt[0]-pt[1]) / 1.414
8         if pt[0] > pt[1]:
9             labels.append(0)
10        else:
11            labels.append(1)
12
13    return np.array(inputs), np.array(labels).reshape(n, 1)
```



- XOR

```
1 def generate_XOR_easy():
2     inputs = []
3     labels = []
4
5     for i in range(11):
6         inputs.append([0.1*i, 0.1*i])
7         labels.append(0)
8
9         if 0.1*i == 0.5:
10            continue
11
12        inputs.append([0.1*i, 1-0.1*i])
13        labels.append(1)
14
15    return np.array(inputs), np.array(labels).reshape(21, 1)
```

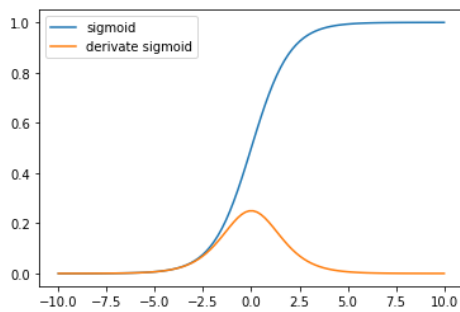


## 2. Experiment setups

### A. Sigmoid functions

```
1 def sigmoid(x):
2     return 1.0 / (1.0 + np.exp(-x))

1 def derivative_sigmoid(x):
2     return np.multiply(x, 1.0 - x)
```



$$\begin{aligned}\sigma(x) &= \frac{1}{1 + e^{-x}} \\ \sigma'(x) &= \frac{d(1 + e^{-x})^{-1}}{dx} \\ &= -(1 + e^{-x})^2 \frac{d}{dx}(1 + e^{-x}) \\ &= -(1 + e^{-x})(1 + e^{-x})(-e^{-x}) \\ &= \sigma(x)(1 - \sigma(x))\end{aligned}$$

Sigmoid是一個映射函數，把變量映射到[0, 1]之間，通常被用來當作neural networks的activation function。

最常使用的情形，就是做binary classification的時候。將model的最後一層layer設定為只有一個neural unit，再將最後輸出的值傳入sigmoid function，就會得到一個介於[0, 1]之間的數值。最後只需要設定threshold，例如將小於0.5的值判斷為0、大於0.5的值判斷為1，就可以做出binary classification的預測。

## B. Neural network

```

1  class myNet():
2      def __init__(self, sizes, learning_rate, activation):
3          self.learning_rate = learning_rate
4          sizes_out = sizes[1:] + [0]
5          self.layer = []
6          for a, b in zip(sizes, sizes_out):
7              if b == 0:
8                  continue
9              elif b == 1:
10                 self.layer += [myLayer(a, b, 'Sigmoid')]
11             else:
12                 self.layer += [myLayer(a, b, activation)]
13
14     def forward(self, x):
15         for l in self.layer:
16             x = l.forward(x)
17         return x
18
19     def backward(self, dC):
20         for l in self.layer[::-1]:
21             dC = l.backward(dC)
22         return dC
23
24     def update(self):
25         gradients = []
26         for l in self.layer:
27             gradients += [l.update(self.learning_rate)]
28         return gradients

```

## C. Backpropagation

- Chain Rule

目標是要minimize loss function的cost C，但 $\partial C / \partial w$ 不易計算，所以要用chain rule

$$\frac{\partial C}{\partial w} = \frac{\partial z}{\partial w} \frac{\partial C}{\partial z}$$

- Forward

```

1  def forward(self, x):
2      self.forward_pass = x
3      self.y = np.matmul(x, self.w)
4      if self.activation == 'Sigmoid':
5          self.y = sigmoid(self.y)
6
7      return self.y

```

算式：  $\frac{\partial z}{\partial w} = \frac{\partial x'w}{\partial w} = x'$

- Backward

```

1  def backward(self, derivative_C):
2      if self.activation == 'Sigmoid':
3          self.backward_pass = np.multiply(
4              derivative_sigmoid(self.y),
5              derivative_C
6          )
7      elif self.activation == 'None':
8          self.backward_pass = derivative_C
9
10     return np.matmul(self.backward_pass, self.w.T)

```

算式：  $\frac{\partial C}{\partial z} = \frac{\partial y}{\partial z} \frac{\partial C}{\partial y}$      $y = \sigma(z), \frac{\partial y}{\partial z} = \sigma'(z)$

- Gradient descent

```

1  def update(self, learning_rate):
2      self.gradient = np.matmul(
3          self.forward_pass.T,
4          self.backward_pass
5      )
6      self.w -= learning_rate * self.gradient
7      return self.gradient

```

算式：  $w = w - \eta \Delta w$

### 3. Result of your testing

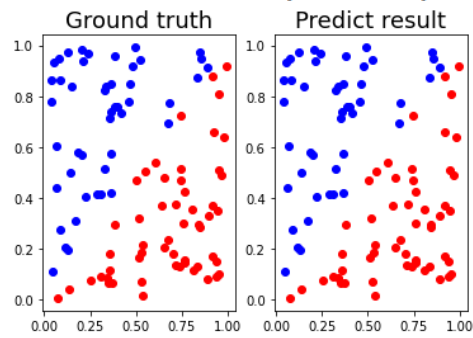
---

With 4 hidden units and sigmoid function.

#### A. Screenshot and comparison figure

- Linear

epoch 500	loss : 0.5085	[[9.99999940e-01]
epoch 1000	loss : 0.2066	[3.72272842e-06]
epoch 1500	loss : 0.1214	[9.99973649e-01]
epoch 2000	loss : 0.0828	[9.99999931e-01]
epoch 2500	loss : 0.0615	[3.51397975e-05]
epoch 3000	loss : 0.0482	[9.99999961e-01]
epoch 3500	loss : 0.0393	[4.96985701e-06]
epoch 4000	loss : 0.0329	[9.99997386e-01]
epoch 4500	loss : 0.0282	[4.35247145e-06]
epoch 5000	loss : 0.0245	[9.99999937e-01]
epoch 5500	loss : 0.0217	[3.70567958e-06]
epoch 6000	loss : 0.0193	[9.99999938e-01]
epoch 6500	loss : 0.0174	[9.99491535e-01]
epoch 7000	loss : 0.0158	[3.61019610e-06]
epoch 7500	loss : 0.0145	[3.65872731e-06]
epoch 8000	loss : 0.0133	[9.99999962e-01]
epoch 8500	loss : 0.0123	[3.87776966e-06]
epoch 9000	loss : 0.0114	[3.57111566e-06]
epoch 9500	loss : 0.0107	[9.99734086e-01]
epoch 10000	loss : 0.0100	[9.99993623e-01]
epoch 10500	loss : 0.0094	[5.08162284e-05]
epoch 11000	loss : 0.0088	[9.99999624e-01]
epoch 11500	loss : 0.0084	[9.99999962e-01]
epoch 12000	loss : 0.0079	[5.90316637e-06]
epoch 12500	loss : 0.0075	[9.99999954e-01]
epoch 13000	loss : 0.0071	[9.99999932e-01]
epoch 13500	loss : 0.0068	[9.999999540e-01]
epoch 14000	loss : 0.0065	[4.44856100e-06]
epoch 14500	loss : 0.0062	[9.99999962e-01]
epoch 15000	loss : 0.0060	[9.99999934e-01]
epoch 15500	loss : 0.0057	[9.99999955e-01]
epoch 16000	loss : 0.0055	[1.78554628e-03]
epoch 16500	loss : 0.0053	[9.99999933e-01]
epoch 17000	loss : 0.0051	[9.99779406e-01]

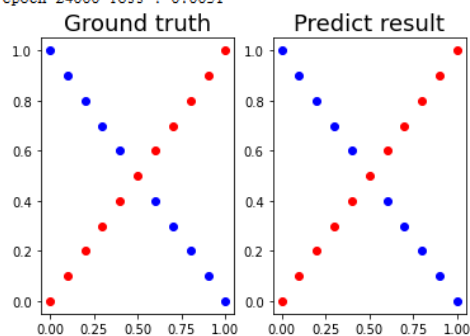


- XOR

```

epoch 1000 loss : 0.3908
epoch 1500 loss : 0.1575
epoch 2000 loss : 0.0966
epoch 2500 loss : 0.0693
epoch 3000 loss : 0.0539
epoch 3500 loss : 0.0441
epoch 4000 loss : 0.0372
epoch 4500 loss : 0.0322
epoch 5000 loss : 0.0284
epoch 5500 loss : 0.0254
epoch 6000 loss : 0.0229
epoch 6500 loss : 0.0209
epoch 7000 loss : 0.0192
epoch 7500 loss : 0.0178 [[3.32711933e-04]
epoch 8000 loss : 0.0165 [9.99673376e-01]
epoch 8500 loss : 0.0155 [3.31094709e-04]
epoch 9000 loss : 0.0145 [9.99672752e-01]
epoch 9500 loss : 0.0137 [3.29511369e-04]
epoch 10000 loss : 0.0129 [9.99675201e-01]
epoch 10500 loss : 0.0123 [3.27961683e-04]
epoch 11000 loss : 0.0116 [9.99672727e-01]
epoch 11500 loss : 0.0111 [3.26445505e-04]
epoch 12000 loss : 0.0106 [9.99077987e-01]
epoch 12500 loss : 0.0101 [3.24962776e-04]
epoch 13000 loss : 0.0097 [3.23513524e-04]
epoch 13500 loss : 0.0093 [9.99139958e-01]
epoch 14000 loss : 0.0090 [3.22097878e-04]
epoch 14500 loss : 0.0087 [9.99752070e-01]
epoch 15000 loss : 0.0084 [3.20716076e-04]
epoch 15500 loss : 0.0081 [9.99768377e-01]
epoch 16000 loss : 0.0078 [3.19368471e-04]
epoch 16500 loss : 0.0075 [9.99771663e-01]
epoch 17000 loss : 0.0073 [3.18055542e-04]
epoch 17500 loss : 0.0071 [9.99771045e-01]]
epoch 18000 loss : 0.0069
epoch 18500 loss : 0.0067
epoch 19000 loss : 0.0065
epoch 19500 loss : 0.0063
epoch 20000 loss : 0.0062
epoch 20500 loss : 0.0060
epoch 21000 loss : 0.0059
epoch 21500 loss : 0.0057
epoch 22000 loss : 0.0056
epoch 22500 loss : 0.0055
epoch 23000 loss : 0.0053
epoch 23500 loss : 0.0052
epoch 24000 loss : 0.0051

```



## B. Show the accuracy of your prediction

```

1 | print(f'Accuracy : {float(np.sum(y == pred_y)) * 100 / len(y)} %')

```

- Linear

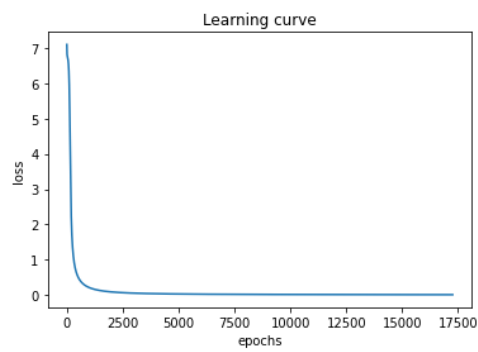
```
Accuracy : 100.0 %
```

- XOR

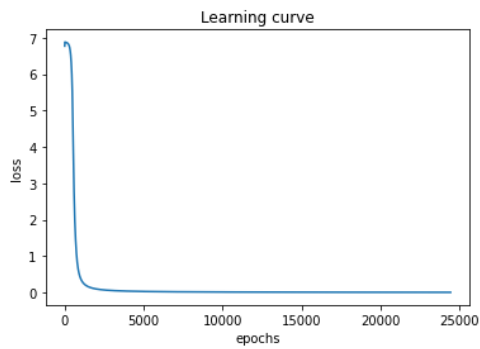
```
Accuracy : 100.0 %
```

## C. Learning curve (loss, epoch curve)

- Linear



- XOR

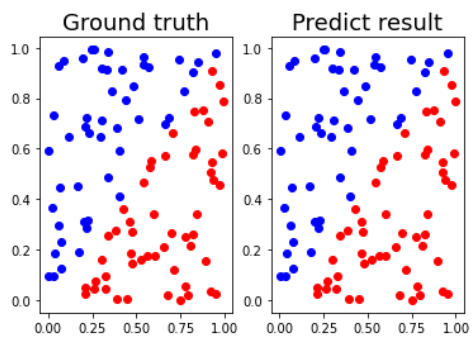
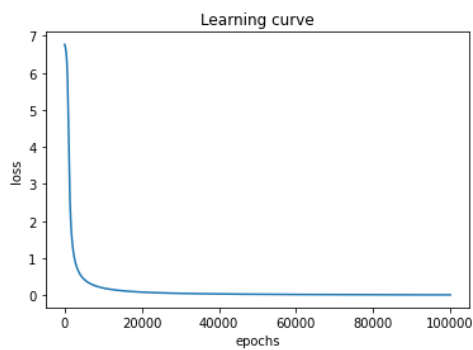


## 4. Discussion

### A. Try different learning rates

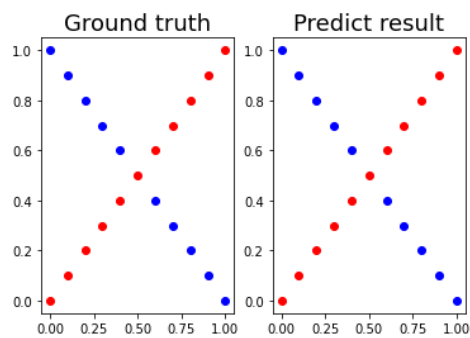
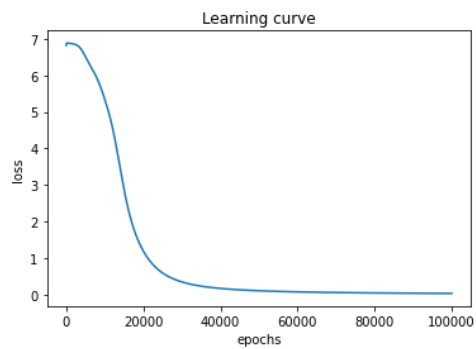
- Learning rate = 0.1

Linear



Accuracy : 100.0 %

## XOR

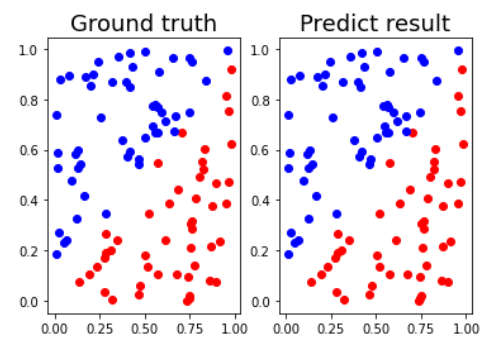
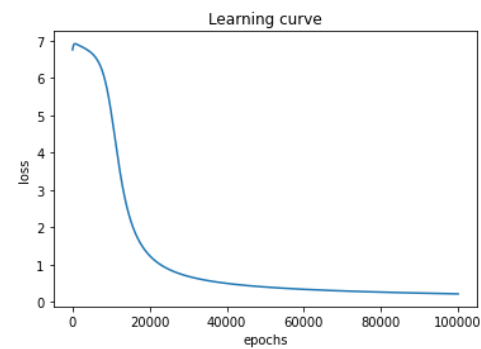


Accuracy : 100.0 %

跟learning rate = 1比起來，loss下降的較慢，在epoch 100000的loss比learning rate = 1的epoch 25000的loss還要高，但accuracy一樣可以達到100%。

- Learning rate = 0.01

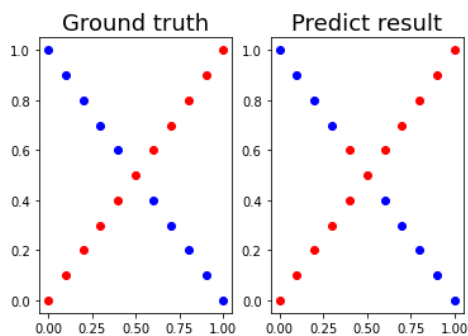
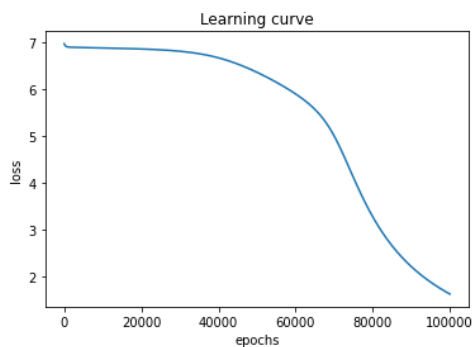
## Linear



Accuracy : 100.0 %



## XOR

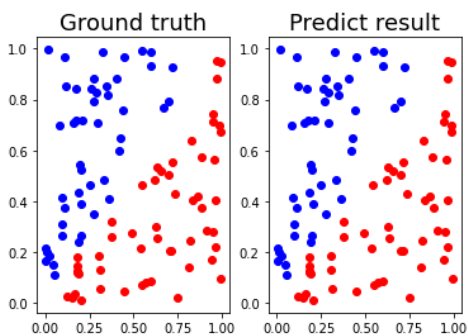
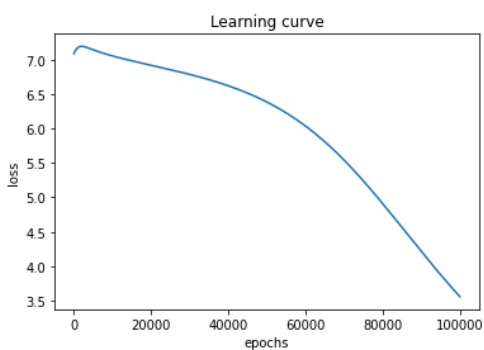


Accuracy : 95.23809523809524 %

learning rate = 0.01時，雖然在learning curve上看起來還是有持續下降的，但非常緩慢，可能需要更多的epoch才會達到更高的accuracy ( 若沒有early stop，最多只會跑100000個epoch )。

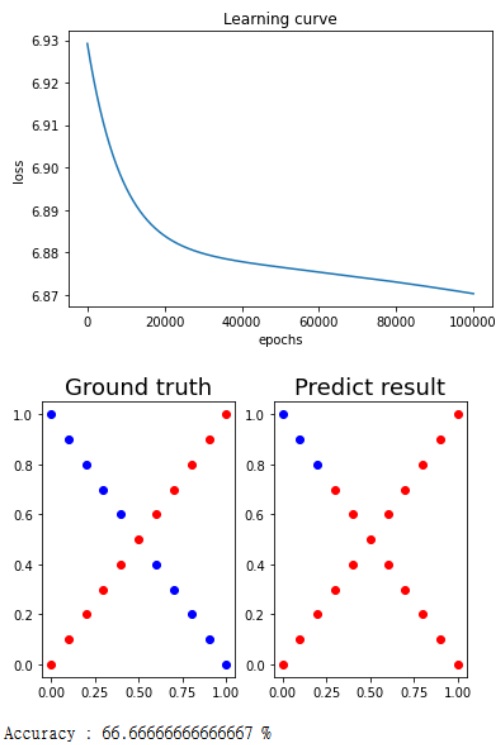
- Learning rate = 0.001

## Linear



Accuracy : 100.0 %

## XOR

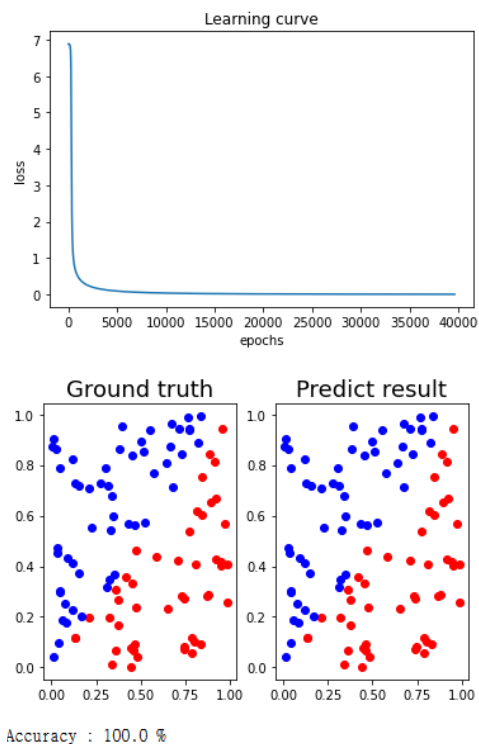


learning rate = 0.001時，loss看起來又下降得更慢了，所以在相同epoch 100000的情況，accuracy就會更低。

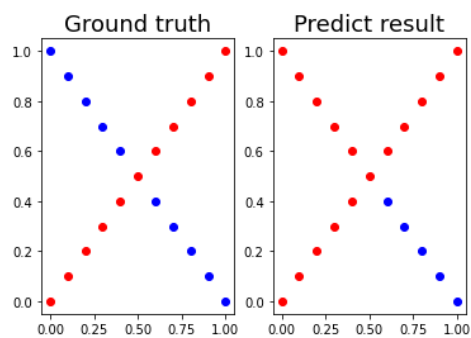
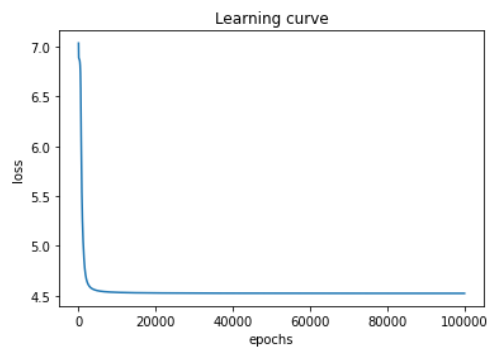
## B. Try different numbers of hidden units

- 2 hidden units

### Linear



## XOR

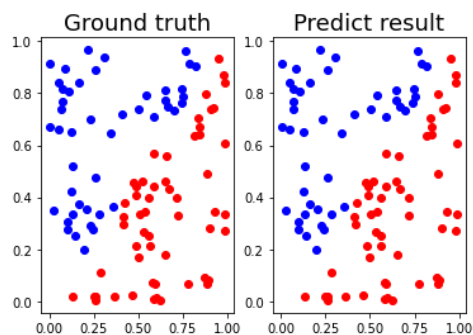
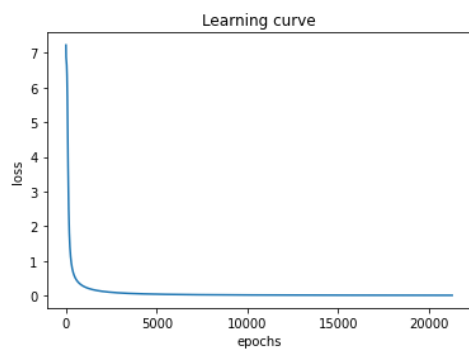


Accuracy : 76.19047619047619 %

相較於4 hidden units · 2 hidden units的loss下降的較慢，並且在XOR data上沒辦法有很好的accuracy。

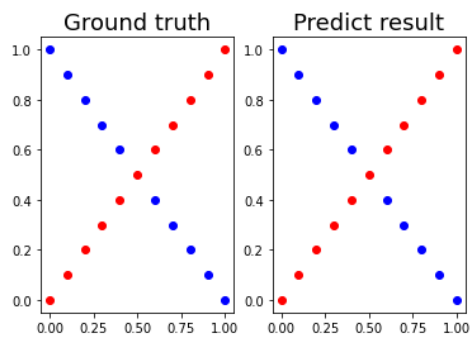
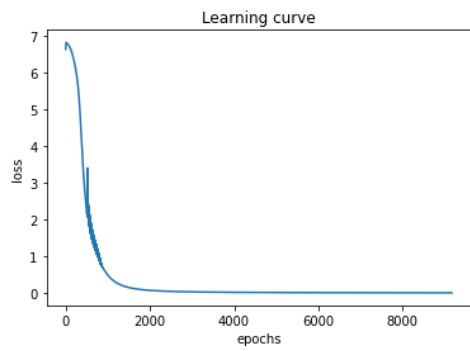
- 8 hidden units

## Linear



Accuracy : 100.0 %

## XOR

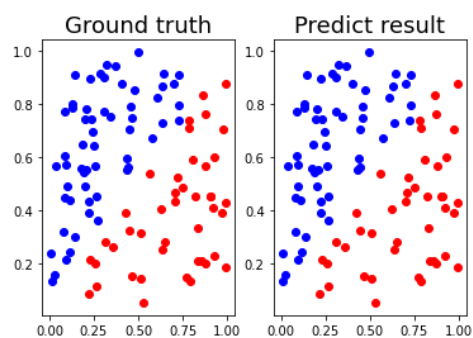
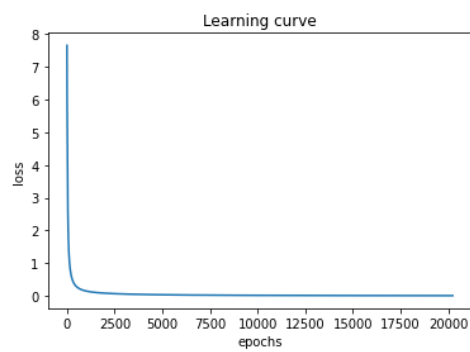


Accuracy : 100.0 %

相較於4 hidden units · 8 hidden units在XOR data的loss收斂得更快。

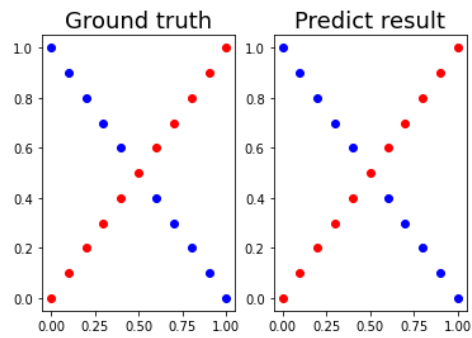
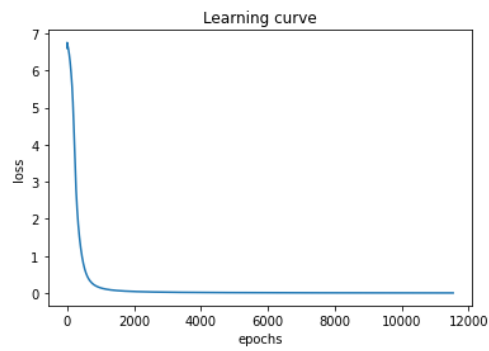
- 16 hidden units

## Linear



Accuracy : 100.0 %

## XOR

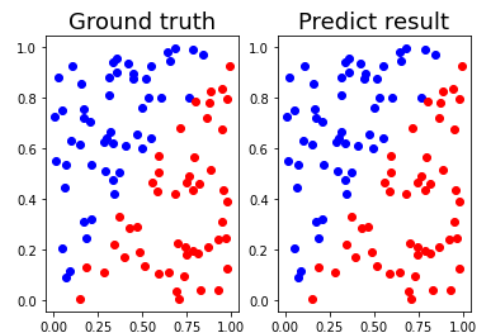
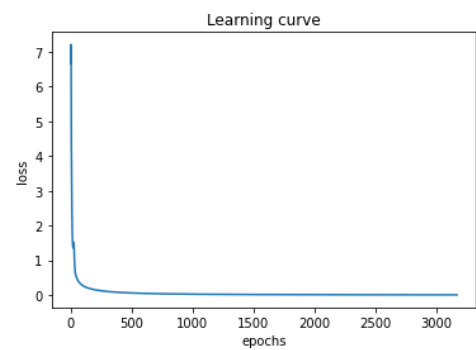


Accuracy : 100.0 %

16和4 hidden units較無差異

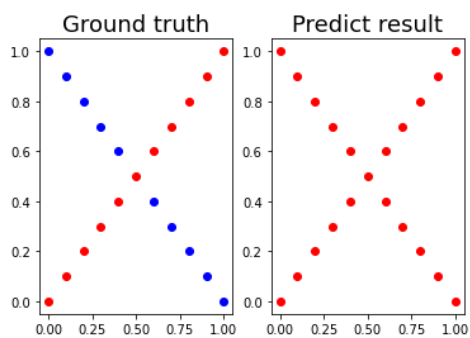
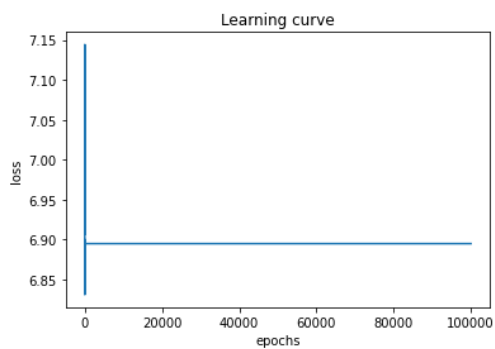
## C. Try without activation functions

- Linear



Accuracy : 100.0 %

- XOR



Accuracy : 52.38095238095238 %

因為XOR data非線性，所以拿掉sigmoid後結果會變得非常差